

Predicting NYPD Response Times after Dispatch

Max Grove | mg6392@nyu.edu

Big Data | Dr. Juan Rodriguez

Tandon School of Engineering

New York University

Outline

- Introduction
 - Data
 - Pipeline of application
 - Loading Data
 - Sampling Data
 - Cleaning Data
 - Exploring Data
 - Modeling
 - Conclusion
-

Introduction

Goal

Predict how long it takes a dispatched NYPD officer to arrive at the scene

DataSet

- **NYC's Open Data Initiative**
- **Documented information on NYPD 911 calls and responses**
- **Data is operated by members of the public *and* NYPD members**



Two datasets:

- **Historic**
 - 2.3 GB size
 - 20 columns
 - 40.7M Rows
- **YTD**
 - ~900 MB size
 - 18 columns
 - 3.6M rows

Operated by members of the public *and* NYPD professionals

Data Definition

Identifiers

objectid	Number
----------	------------------------

cad_evnt_id	Number
-------------	------------------------

Miscellaneous

nypd_pct_cd	Number
-------------	------------------------

radio_code	Text
------------	----------------------

typ_desc	Text
----------	----------------------

cip_jobs	Text
----------	----------------------

Timestamps

create_date	Floating Timestamp
-------------	------------------------------------

incident_date	Floating Timestamp
---------------	------------------------------------

incident_time	Text
---------------	----------------------

add_ts	Floating Timestamp
--------	------------------------------------

disp_ts	Floating Timestamp
---------	------------------------------------

arrivd_ts	Floating Timestamp
-----------	------------------------------------

closng_ts	Floating Timestamp
-----------	------------------------------------

Location

boro_nm	Text
---------	----------------------

patrl_boro_nm	Text
---------------	----------------------

geo_cd_x	Number
----------	------------------------

geo_cd_y	Number
----------	------------------------

latitude	Number
----------	------------------------

longitude	Number
-----------	------------------------

location	Point
----------	-----------------------

Introduction

Data

Pipeline

Loading Data

Sampling
Data

Cleaning
Data

Evaluating
Data

Model

Conclusion

Pipelining Data and Technology



Collecting and understanding Data



Hosting data and indexing



Cleaning data
Exploring data
Modeling data

Introduction

Data

Pipeline

Loading Data

Sampling
Data

Cleaning
Data

Evaluating
Data


Model

Conclusion

Loading Data to MongoDB

```
api_endpoint_ytd = "https://data.cityofnewyork.us/resource/n2zq-pubd.json"
```

```
api_endpoint_historic = "https://data.cityofnewyork.us/resource/d6zx-ckhd.json"
```



- Create an App Token with NYC Open Data
- Fetch the data from the APIs - 20,000 data points per call
- Insert into a Mongo DB collection

Key MongoDB Features Used



Introduction

Data

Pipeline

Loading Data

Sampling
Data

Cleaning
Data


Evaluating
Data

Model

Conclusion

Sampling Data from MongoDB

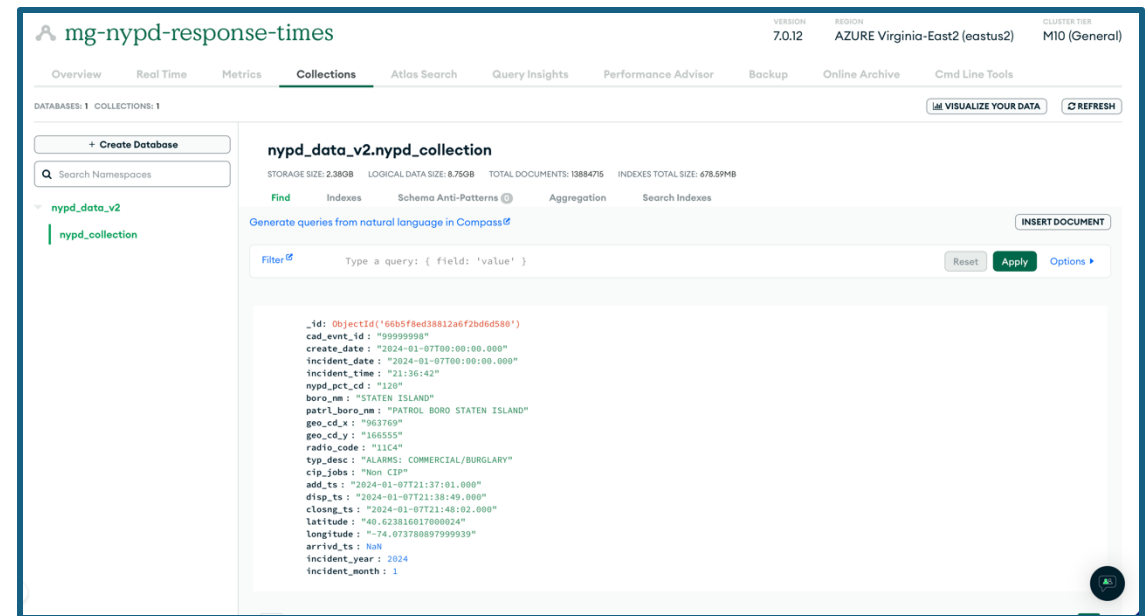
Pull all available month-year combinations from MongoDB



Pull 1,000 rows from each month-year down into the local JupyterHub environment

Key Benefits

- Save on space
- Ensure we get data from all time periods
- Utilize indexed rows
- Don't have to rely on API and data is backed up



Cleaning and Preparing Data

Initial Schema

22 columns and 22,000 rows

_Id	Typ_desc
Cad_evnt_id	Cip_jobs
Create_date	Add_ts
Incident_date	Disp_ts
Incident_time	Closng_ts
Nypd_pct_cd	Latitude
Boro_nm	Longitude
Patrl_boro_nm	Location
Geo_cd_x	Incident_year
Geo_cd_y	Incident_month
Radio_code	Arrivd_ts

Cleaning

- Dropping unnecessary columns
- Dropping missing values
- Cleaning dates

New Features

- Tokenize words and create feature flags
- New time features
- Linking data to NYC Neighborhood data
- Key data: Time to arrive from dispatch time

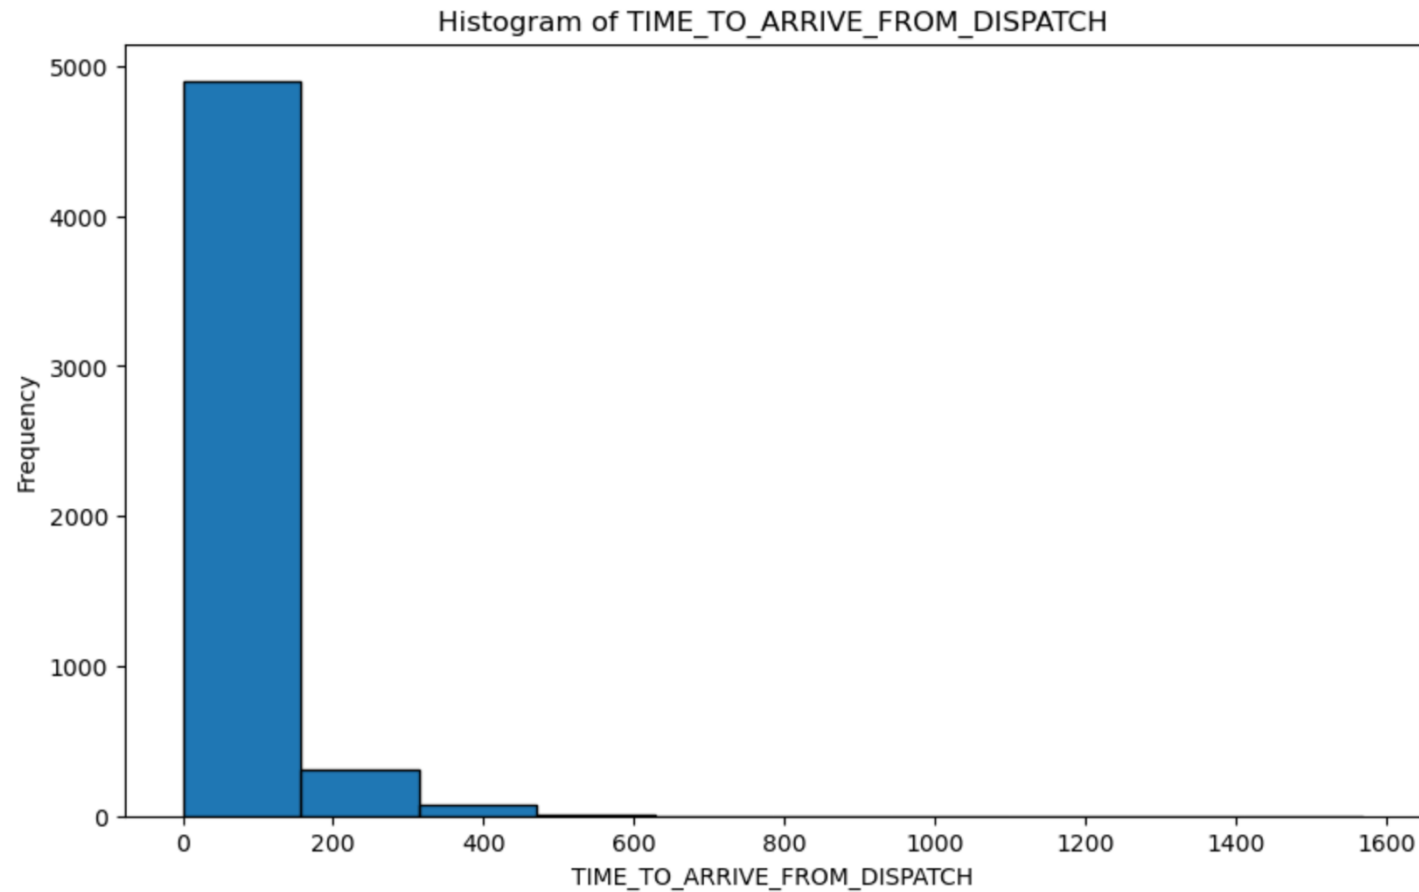
Cleaned Schema

22 columns and 5,291 rows

NYPD_PCT_CD
BORO_NM
RADIO_CODE
CIP_JOBS
INCIDENT_YEAR
INCIDENT_MONTH
TYP_DESC_HAS_{SELECTED_WORDS}
HOUR
WEEKDAY
TIME_TO_ARRIVE_FROM_DISPATCH
NEIGHBORHOOD

* New Features

Evaluating and Understanding Data



Introduction

Data

Pipeline

Loading Data

Sampling
Data

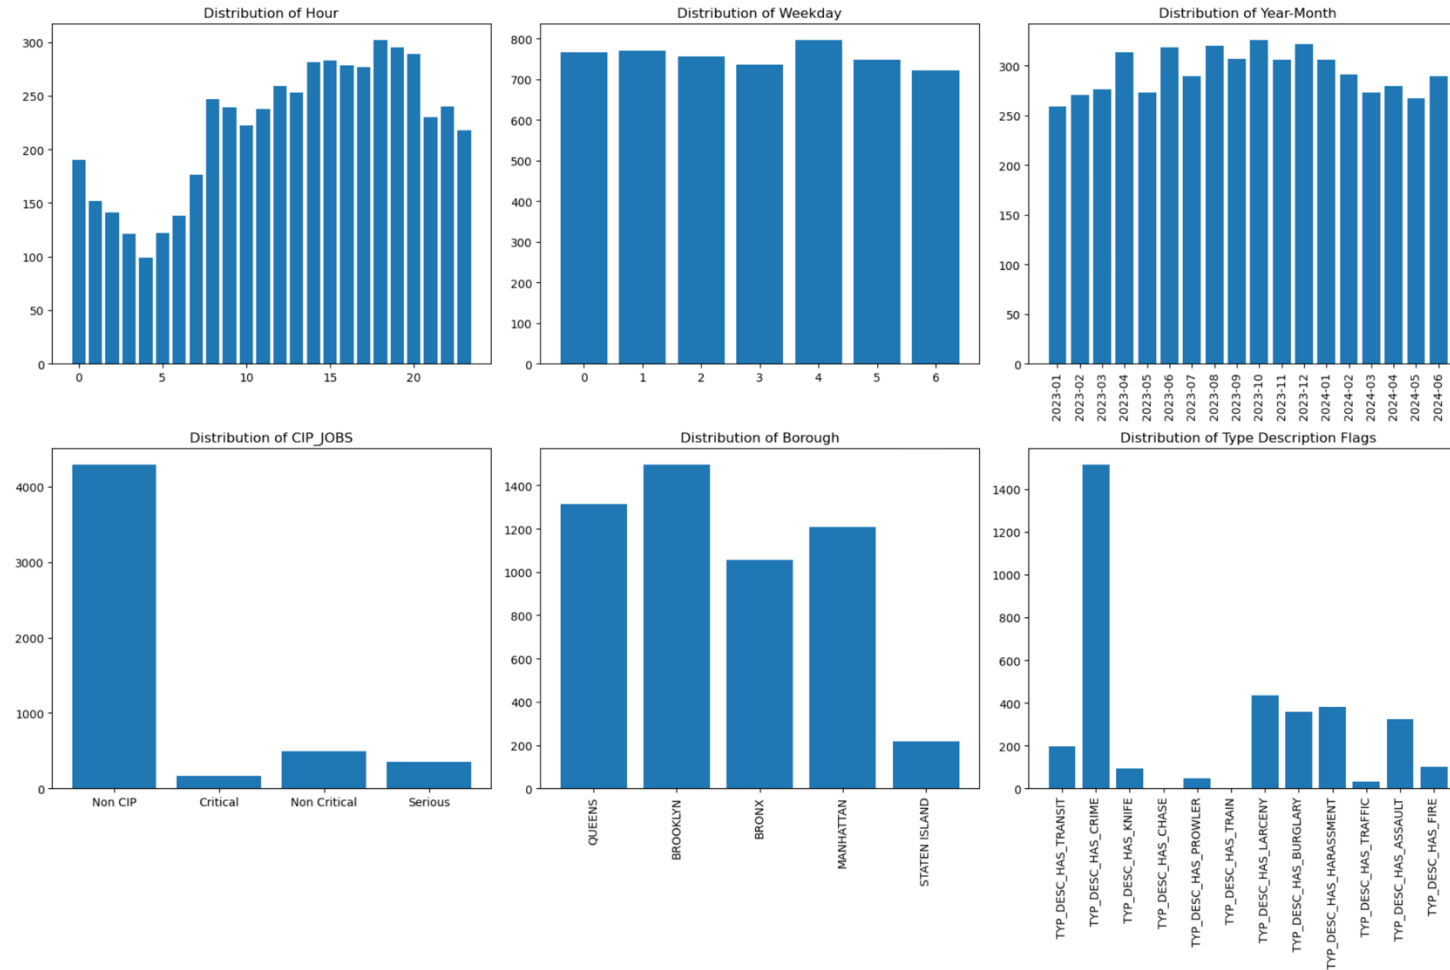
Cleaning
Data

Evaluating
Data

Model

Conclusion

Evaluating and Understanding Data



Introduction

Data

Pipeline

Loading Data

Sampling
Data

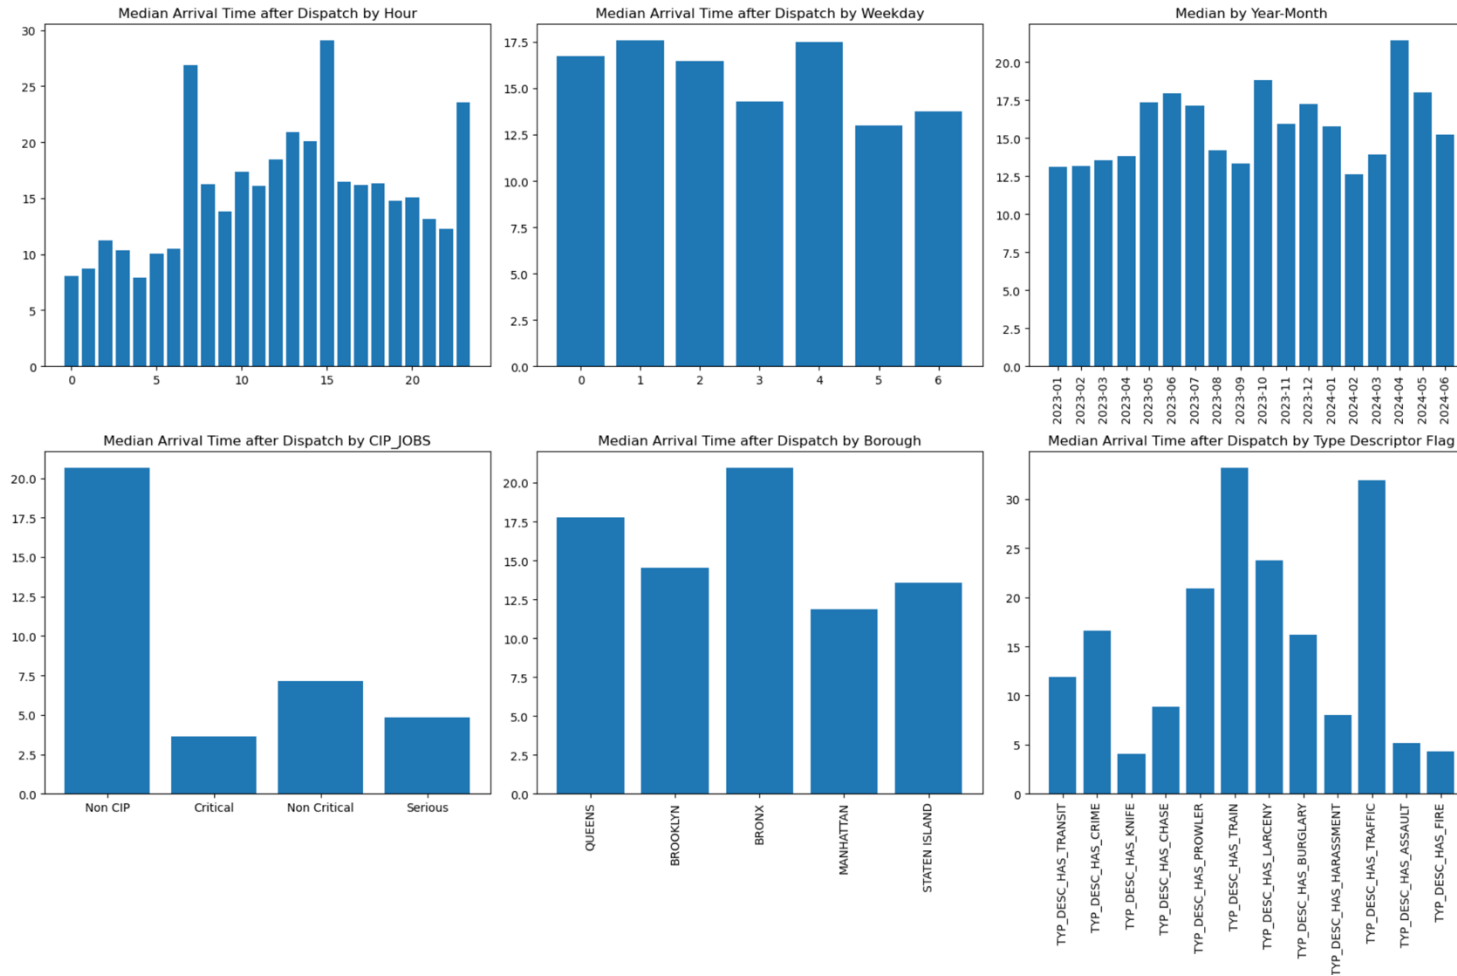
Cleaning
Data

Evaluating
Data

Model

Conclusion

Evaluating and Understanding Data



Introduction

Data

Pipeline

Loading Data

Sampling
Data


Cleaning
Data

Evaluating
Data

Model


Conclusion

Evaluating and Understanding Data



- Red: 100th Percentile
- Orange: 80th Percentile
- Yellow: 60th Percentile
- Yellow-Green: 40th Percentile
- Green: Bottom 20th Percentile

Evaluating and Understanding Data



- Red: 100th Percentile
- Orange: 80th Percentile
- Yellow: 60th Percentile
- Yellow-Green: 40th Percentile
- Green: Bottom 20th Percentile

Machine Learning Models

Filter out actual arrival times greater than 40 minutes and less than 3 minutes

```
columns = [  
    "BORO_NM",  
    "RADIO_CODE",  
    "CIP_JOBS",  
    "NEIGHBORHOOD"  
]
```

Create index
values off of String
columns

Create a test and training data set

Training data set: 80%
Test Data Set: 20%

Introduction

Data

Pipeline

Loading Data

Sampling
Data


Cleaning
Data

Evaluating
Data

Model

Conclusion

Machine Learning Models



RMSE: 9.27

Introduction

Data

Pipeline

Loading Data

Sampling
Data

Cleaning
Data

Evaluating
Data

Model

Conclusion


Machine Learning Models

Rerun model with tuning

- Parameter Grid
- Cross Validator with 5 folds

Hyperparameter Takeaways

- Hyperparameter tuning can lead to overfitting on the training data
- Chosen parameter grid might not be optimal
- Can sometimes make the model too restrictive or too flexible



RMSE: 9.31

Introduction

Data

Pipeline

Loading Data

Sampling
Data

Cleaning
Data

Evaluating
Data

Model

Conclusion

Conclusion

Takeaways

- Can be incredibly difficult to clean and parse data from public sources
- Models often need lots of work to get them in shape
- Can be incredibly difficult to predict things that have many more inputs than what can be represented in data

Next Steps

- Further tune the model by removing ancillary features
- Run the model Borough by borough, or even neighborhood by neighborhood
- Run the model through Dask or other ML Models
- Bring in more historic data
- Bring in weather data to see if weather on a given day affects the arrival time